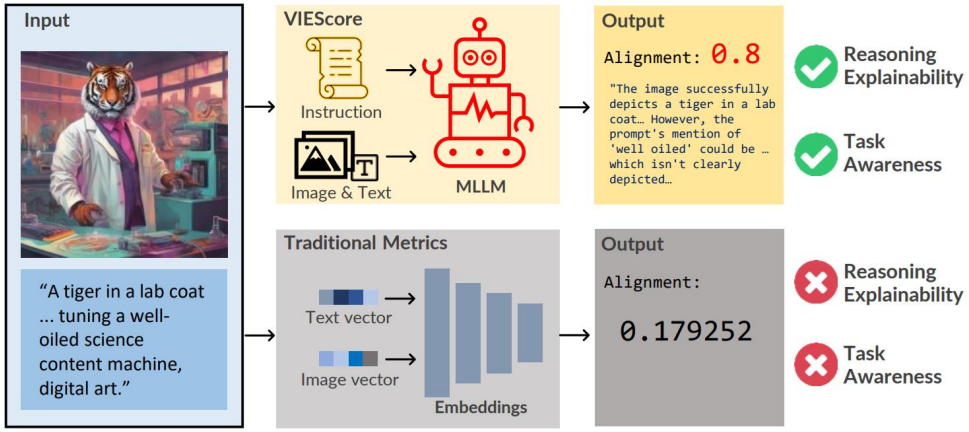

VIEScore: Towards Explainable Metrics for Conditional Image Synthesis Evaluation

♠Max Ku, ♠Dongfu Jiang, ♠Cong Wei, ♥Xiang Yue, ♠Wenhu Chen

♠University of Waterloo, ♥IN.AI Research

tiger-ai-lab.github.io/VIEScore

Paper ID: #2258

ACL 2024 Bangkok, Thailand

UNIVERSITY OF WATERLOO

Motivation

Metrics in the future would provide not just the score but also the rationale

- Understanding of judgment per instance is needed
- Traditional metrics are not task aware

Visual Instruction-guided Explainable Score (VIEScore)

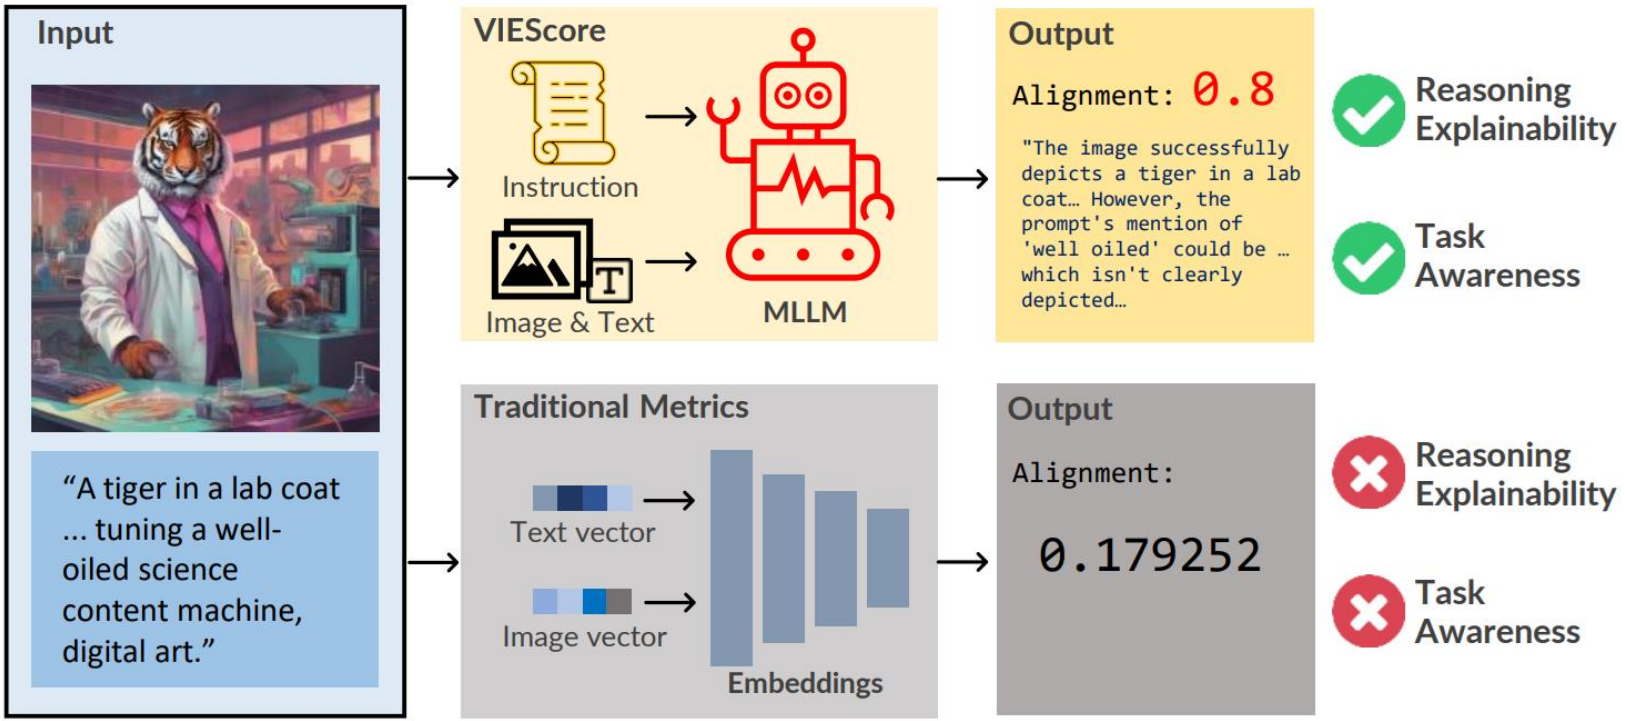

How VIEScore works?

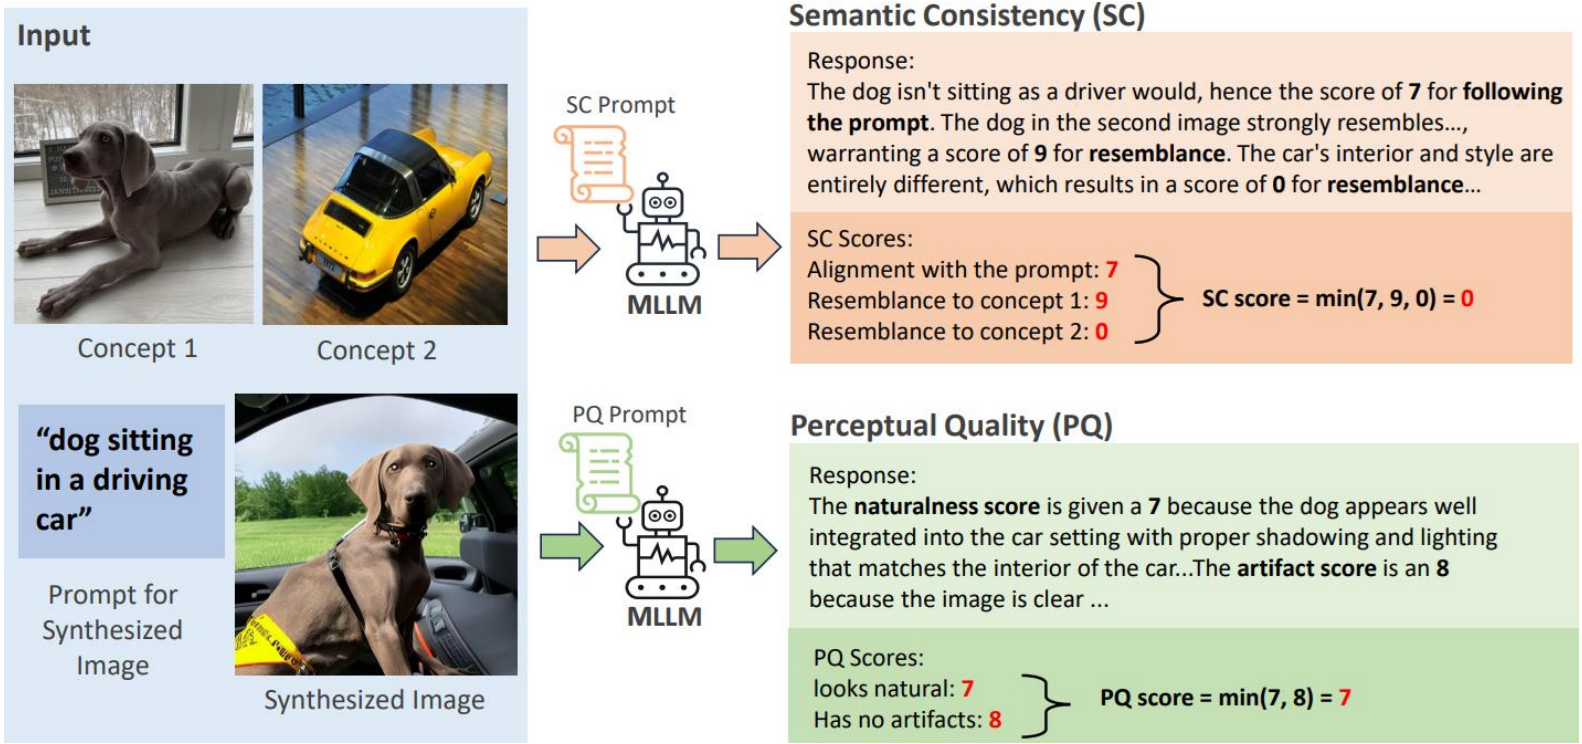

Experiment Setup

A wide range of image synthesis tasks study

- Correlation of VIEScore to Human
- V.S.
- Correlation of Traditional metrics to Human

Where can we get this kind of human annotation data?

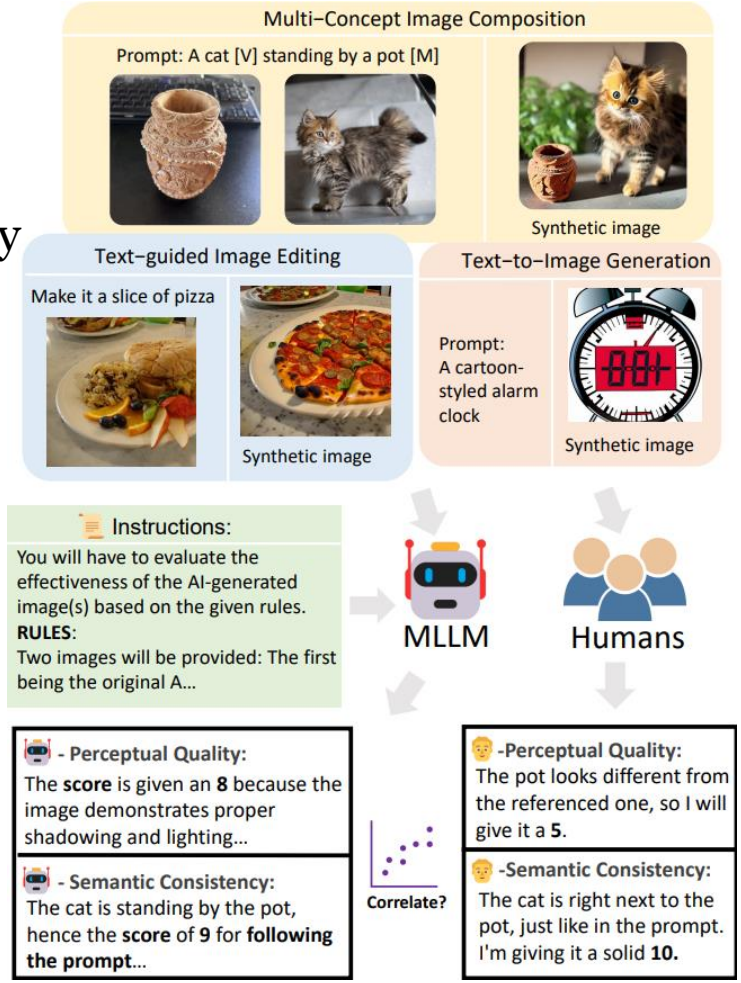

Experiment Setup (Cont.)

- Human data from ImagenHub(ICLR 2024)
- 29 Models across 7 tasks
  - Total 14403 annoations
- Each annotation 3 human metrics:
  - SC : Conditions-Image alignment
  - PQ: Realism and Natural sense
  - Overall: sqrt(SC x PQ)
- Human guideline is used as prompt

c_1	c_2	c_2	Task	y
A cartoon styled alarm clock	∅	∅	Text-to-Image Generation	
		Change frisbee to a football	Mask-guided Image Editing	
	Make it a slice of pizza instead of the sandwich	∅	Text-guided Image Editing	
	A [V] dog in the Versailles hall of mirrors	∅	Subject-Driven Image Generation	
		Replace glasses with [V] glasses	Subject-Driven Image Editing	
		A cat [V] standing by a pot [M]	Multi-Concept Image Composition	
	A small dog is curled up on top of the shoes	∅	Control-guided Image Generation	

ImagenHub: Standardizing the evaluation of conditional image generation models
https://arxiv.org/abs/2310.01596

Main Result

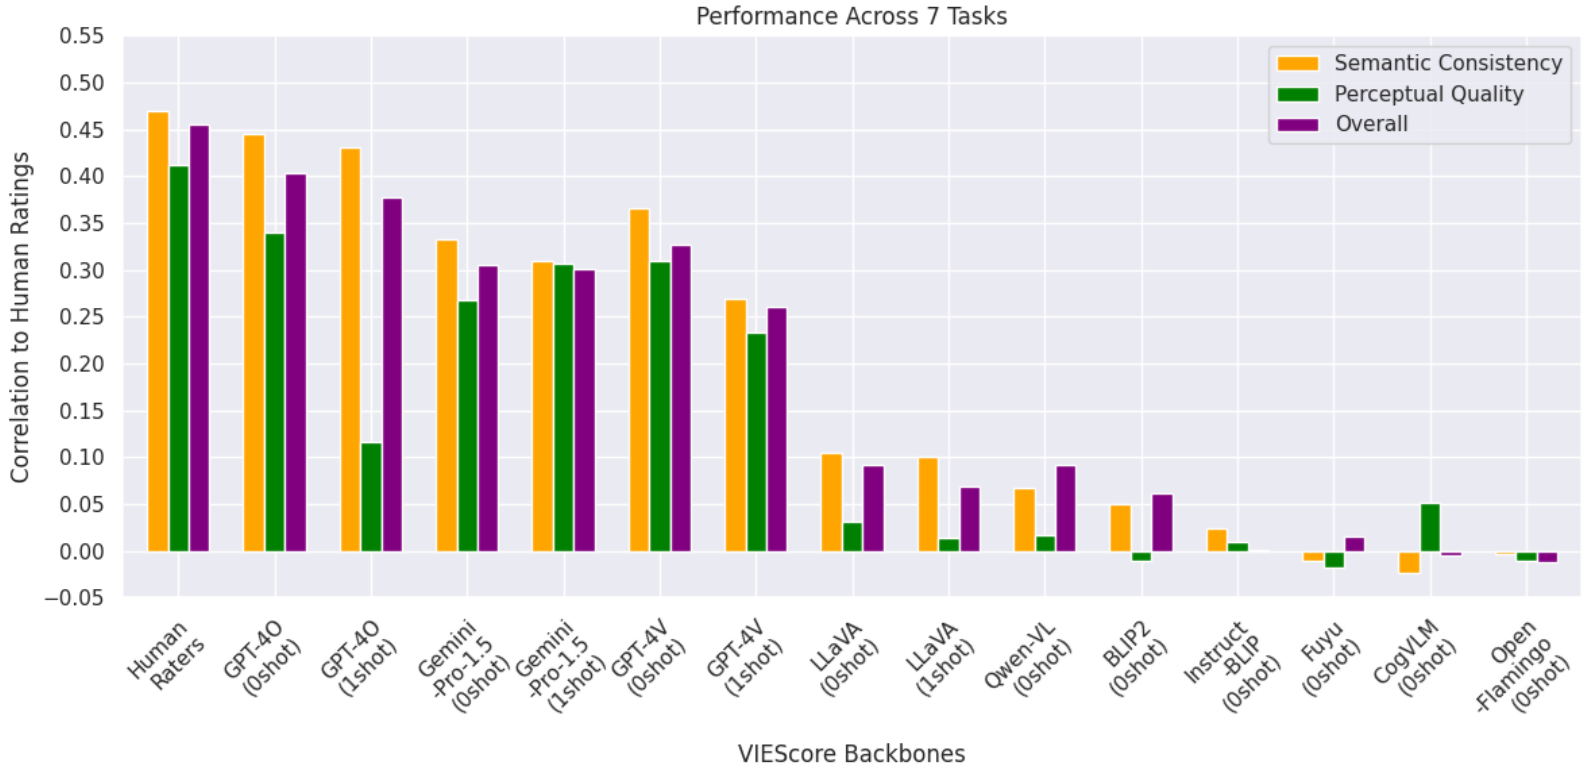
Performance Across 7 Tasks
Semantic Consistency
Perceptual Quality
Overall
Correlation to Human Ratings
0.55
0.50
0.45
0.40
0.35
0.30
0.25
0.20
0.15
0.10
0.05
0.00
−0.05
Human Raters
GPT-4O (0shot)
GPT-4O (1shot)
Gemini -Pro-1.5 (0shot)
Gemini -Pro-1.5 (1shot)
GPT-4V (0shot)
GPT-4V (1shot)
LLaVA (0shot)
LLaVA (1shot)
Qwen-VL (0shot)
BLIP2 (0shot)
Instruct -BLIP (0shot)
Fuyu (0shot)
CogVLM (0shot)
Open -Flamingo (0shot)
VIEScore Backbones

Why one-shot setting achieve worse performance?

- MLLMs struggle in In-Context Learning when multiple images exists
  - Reasoning is affected

- Appears on all MLLMs we benchmarked

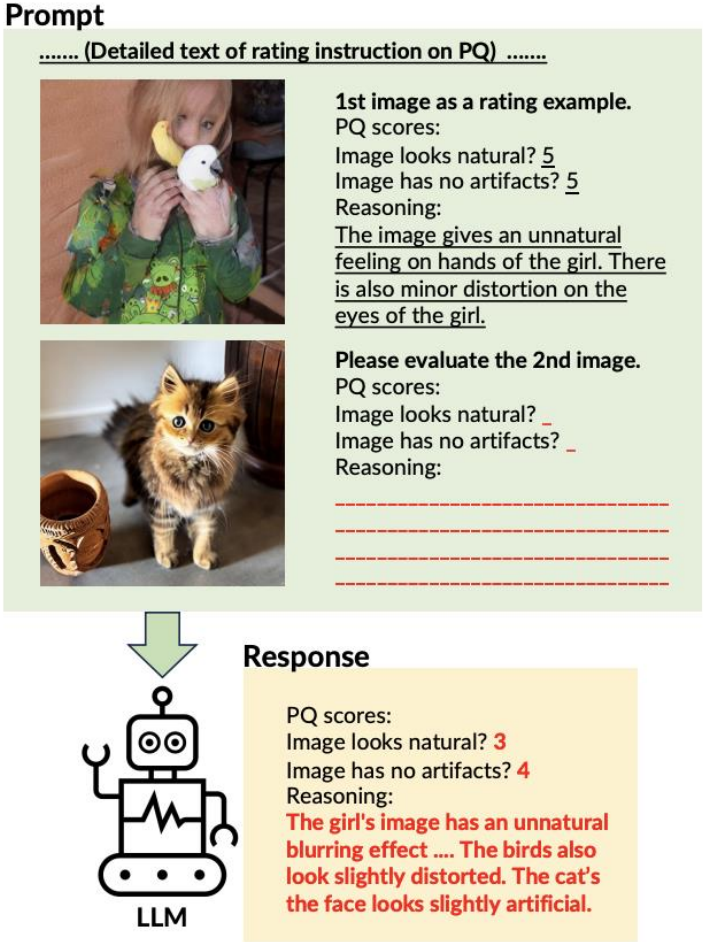

VIEScore achieve higher correlation than Traditional metrics

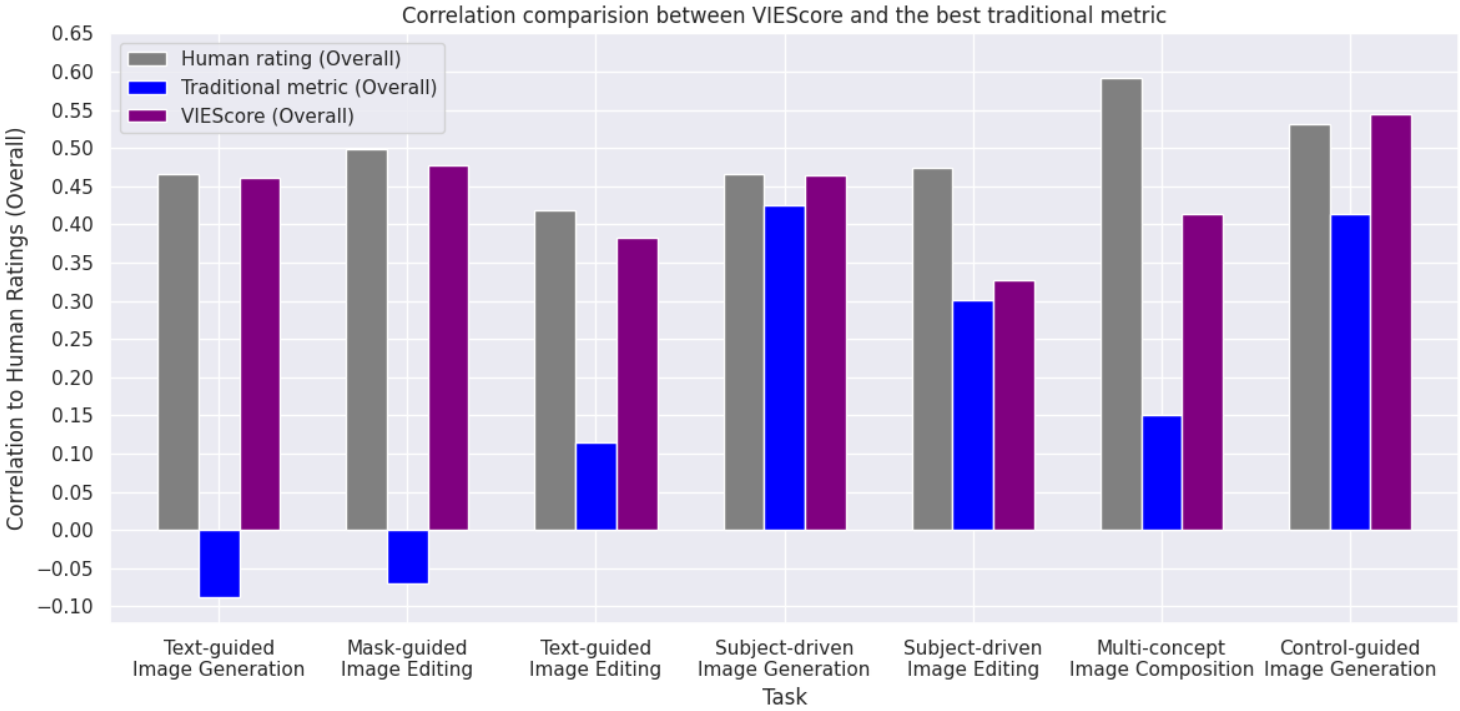

There’s still obstacles toward explainable metrics

- MLLMs struggles to see minor difference when given 2 similar images

Code Release

- Code available on GitHub
- Easy to use and modify

```
from viescore import VIEScore
backbone = "gemini"
vie_score = VIEScore(backbone=backbone, task="t2v")

score_list = vie_score.evaluate(pil_image, text_prompt)
sementics_score, quality_score, overall_score = score_list
```